NEBIUS Ecosystem recommendation: three partners, two quarters, four FDEs

To the Nebius exec team · from Kwasi Ankomah · Nov 2026 · one page

Recommendation. Over the next two quarters, land **Langfuse** (observability), **Cartesia** (voice) and **Daytona** (agent sandboxes), in that order, and staff four FDE hires so the same kit lands ten partners by the end of Q2. The choice is deliberate: one partner proves where Token Factory's premium inference wins (voice), one gives us distribution into coding agents (sandboxes), and one instruments everything so field learning becomes post-training data on our own platform (observability). None duplicates something we already bought (Tavily, ConTree, Clarifai). The repo that accompanies this page already runs against all three (Part 1): a Daytona sandbox proves the framework's tool call, Cartesia speaks the verdict, Langfuse receives every trace and the app reads its metrics back.

The three tests we applied

1. **Premium tokens.** Does the partner put Token Factory in a path where latency *is* the product?
2. **Reach we don't have.** Does it bring developers who would not otherwise pick Nebius?
3. **The loop.** Does it feed traces → evals → datasets → post-training on Token Factory?

The partners

#	Partner	Why now	What Nebius offers	What we get	Lands
1	**Langfuse** — open-source LLM observability, OTel-native, self-hostable	Cheapest, fastest, and it instruments partners 2 and 3. Every reference agent we ship already emits OTel (see repo).	"Langfuse on Nebius": a one-command self-host on Managed Kubernetes next to the inference it traces; Token Factory as a default model in Langfuse playground/evals; Builder Program credits and co-marketing	Traces on partner traffic → eval datasets → post-training with the Eigen team. The flywheel that makes open models on Nebius better than the same weights elsewhere	Q1, weeks 1–6
2	**Cartesia** — real-time voice (Sonic TTS, Ink STT, Line agent platform)	Voice has a ~200 ms conversational budget; the LLM gets a slice. It is the clearest case where fast inference wins, and where we have to prove it on their traffic, not a leaderboard	Dedicated Token Factory endpoints with a TTFT SLA for the LLM leg (Inferize scaling absorbs call spikes); Nebius GPU capacity to host Cartesia's own models; a joint "voice agent on Nebius" reference	Premium real-time tokens and the public proof point that Token Factory is the inference for voice. Hume on SambaNova is the precedent	Start Q1 week 3, live by end Q2
3	**Daytona** — open-source sandboxes for agents	Coding agents are loops: every token is many tokens. Daytona has the developer base; we have the fast decode	Daytona's sandbox fleet on Nebius compute; Token Factory as the default model in Daytona's agent templates; joint coding-agent reference (this repo's `run_python` already supports Daytona and ConTree)	Coding-agent token volume and a route into OSS harnesses (OpenHands, OpenCode) that already use Daytona	Q2

Why this order. Langfuse first because it takes weeks and instruments everything after it. Cartesia starts almost immediately because it is the longest cycle and the biggest revenue proof. Daytona in Q2, after we have agreed how it sits next to ConTree and Token Factory's own Sandboxes beta: the beta runs one fresh operation per call, Daytona gives coding agents the long-lived box, file round-trips, git/LSP and snapshots they need today, and the sandbox interface in the repo takes a Nebius backend when the beta is ready.

Four hires, and how they get us to ten

Hire	Owns	Lands	Sets up
FDE-1 Agents & frameworks	the conformance kit and reference agents; frameworks #2–#5 (LangGraph tests, CrewAI, OpenAI Agents SDK, Mastra)	Daytona (with FDE-4)	Tier-2 self-serve: partners run `verify`, we review
FDE-2 Real-time & voice	TTFT engineering, dedicated endpoints, work with the Inferize team	Cartesia	the premium real-time tier playbook for the next voice partner
FDE-3 Observability & evals	Langfuse on Nebius, OTel conventions, traces → datasets → post-training with Eigen	Langfuse	the loop every later partner plugs into
FDE-4 Platform & developer relations	partner tiering, Builder Program office hours, docs generated from the kit, reference repos	Daytona (with FDE-1)	the machinery that makes partners 4–10 a kit-and-office-hours job, not an engineering project

Tiering is the scaling rule: Tier 1 gets an FDE (the three above); Tier 2 gets the kit, a review and office hours; Tier 3 gets docs generated from the kit's reports. Ten partners by end of Q2 means three Tier 1, four Tier 2, three Tier 3.

The trade-off we are making

We are choosing **premium-inference proof and developer distribution over enterprise breadth.** No enterprise ISV (Palantir-style) and no gateway (OpenRouter, LiteLLM) in the first three; those stay with core engineering and sales for two quarters. And we are choosing Daytona's distribution over exclusivity for ConTree: ConTree stays the first-party default in our docs; Daytona is a supported backend for developers already there.

One concrete ask

A standing partner capacity reserve owned by FDE: two dedicated Token Factory endpoints and a $150k credits pool, so a partner integration never waits on a sales order, and a named Token Factory product counterpart who takes the kit's FAIL reports as roadmap input.

What the live run already tells us (2 Nov 2026)

Token Factory serves 30 models; `Qwen3-235B-A22B-Instruct-2507` passes all 13 kit checks and streams at ~190 tok/s with 80 % prefix-cache hits on an agent loop. Two of the newest arrivals are not agent-ready yet (Kimi-K3 fails `tool_choice` and `json_schema`; DeepSeek-V4.1-Flash ignores `json_schema`), MiniMax-M3.1-Flash has a 3.2 s TTFT on the shared endpoint, and no model meets a 300 ms voice budget there. That is the Cartesia case for dedicated endpoints and the Langfuse case for making these numbers a table every partner can read.

How we will know it worked (end of Q2)

Three partner references live with public READMEs and conformance reports · conformance reports for the top ten served models · Cartesia p50 TTFT under 300 ms on a dedicated endpoint · share of partner traffic traced into Langfuse · tokens per month arriving through partner integrations, reported monthly.